

# Relevance Feedback (and Pseudo Relevance Feedback)

## COMP3009J: Information Retrieval

Dr. David Lillis (david.lillis@ucd.ie)

UCD School of Computer Science  
Beijing Dublin International College

# Relevance Feedback

- **Relevance Feedback** is when the user is involved in the retrieval process to improve the final result set.
- Specifically the user gives feedback on the **relevance** of documents in the initial set of results, and this is used to re-process their query and provide an updated set of results.
- Although this is not typically used in most modern IR systems, the process is important because **pseudo relevance feedback** is common, which is a simulation of relevance feedback without the user being directly involved.

# Relevance Feedback Process

## ■ Typical process:

1. User **provides a query** (usually short and simple)
2. The system returns an **initial set of results**.
3. The user marks some **relevant** and **nonrelevant** documents (in some situations the user only marks relevant results).
4. The system computes a **better representation of the information need** based on this feedback.
5. The system returns a **revised set of results**.
6. Possibly the above process is repeated.

# Rocchio Algorithm: Background

- Probably the best-known algorithm for relevance feedback is the **Rocchio Algorithm**.
- Basic Idea:
  - Based on the idea that documents and queries are **represented as vectors** (e.g. using TF-IDF, although any vector representation will work).
  - Wants to **find a query vector** that maximises similarity with relevant documents and minimises similarity with nonrelevant documents.


# Rocchio Algorithm: Background

- Formally the aim is to find an optimal query vector  $\vec{q}_{opt}$ .
- The optimal query vector is the query vector with the largest difference between its similarity with the set of relevant documents and its similarity with the set of nonrelevant documents.

$$\vec{q}_{opt} = \arg \max_{\vec{q}} [sim(\vec{q}, C_r) - sim(\vec{q}, C_{nr})]$$

- where:
  - $sim(\dots)$  is a measure of similarity (e.g. cosine similarity)
  - $C_r$  is the set of relevant documents
  - $C_{nr}$  is the set of nonrelevant documents

# Rocchio Algorithm: Background



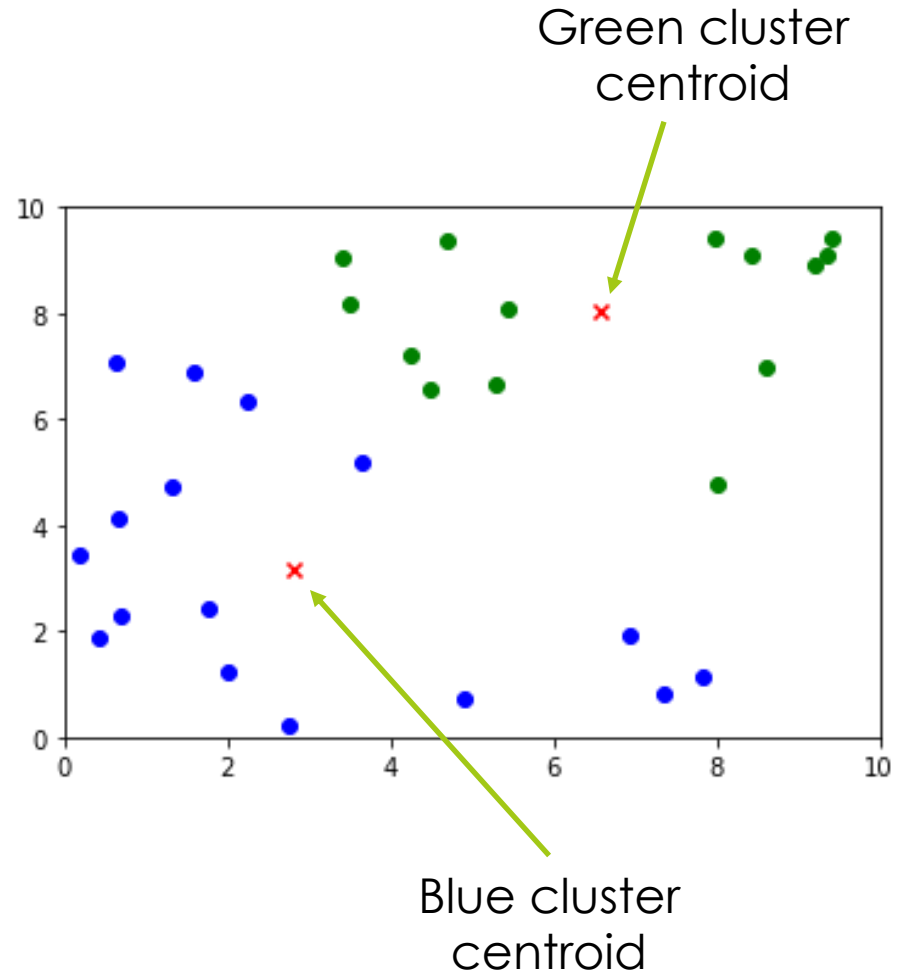
- This illustrates the ideal situation.
- Of course, the real vector space has far more than 2 dimensions: this is only an illustration.

# Centroids

To further examine how the similarity between a query vector and a set of documents can be calculated, we need to consider the idea of a **centroid** of a set of vectors.

The centroid is the centre point of a set of vectors. It is itself a vector.

$$\textit{Centroid}(C) = \frac{1}{|C|} \sum_{\vec{d}_j \in C} \vec{d}_j$$



# Rocchio Algorithm: Theory

## With complete knowledge

- Using cosine similarity, the formula for calculating the optimal query vector is:

$$\vec{q}_{opt} = \frac{1}{|C_r|} \sum_{\vec{d}_j \in C_r} \vec{d}_j - \frac{1}{|C_{nr}|} \sum_{\vec{d}_j \in C_{nr}} \vec{d}_j$$


- Informally, this is the vector that is obtained by subtracting the centroid of the nonrelevant set from the centroid of the relevant set.
- **But:** In practice we don't yet know the relevant and nonrelevant sets!



# Rocchio Algorithm in practice

In practice, in relevance feedback situations, we have access to a subset of the relevant and nonrelevant sets, which have been identified by the user.

The Rocchio algorithm takes this information and computes a **modified query vector** (or “revised query vector”), which is probably not be the optimal query vector, and re-runs retrieval.



# Rocchio Algorithm: Formula

- The formula to calculate the **modified query vector** ( $\vec{q}_m$ ) is as follows:

$$\vec{q}_m = \alpha \vec{q}_0 + \beta \frac{1}{|D_r|} \sum_{\vec{d}_j \in D_r} \vec{d}_j - \gamma \frac{1}{|D_{nr}|} \sum_{\vec{d}_j \in D_{nr}} \vec{d}_j$$

- Here:
- $\vec{q}_0$  is the original query provided by the user.
- $D_r$  and  $D_{nr}$  are the documents that have been identified by the user as being relevant and nonrelevant respectively.
- $\alpha$ ,  $\beta$  and  $\gamma$  are weights that can affect the behaviour of the formula.

# Rocchio Algorithm: Weights

$$\vec{q}_m = \alpha \vec{q}_0 + \beta \frac{1}{|D_r|} \sum_{\vec{d}_j \in D_r} \vec{d}_j - \gamma \frac{1}{|D_{nr}|} \sum_{\vec{d}_j \in D_{nr}} \vec{d}_j$$

- From this formula, what is the effect of changing the weights  $\alpha$ ,  $\beta$  and  $\gamma$ ?
  - They control the balance between trusting the judged sets of relevant and nonrelevant documents and trusting the original query.
  - With many judged documents,  $\beta$  and  $\gamma$  might be high.
  - The modified query moves the original query some distance away from the centroid of the nonrelevant documents ( $\gamma$ ) and some distance towards the centroid of the relevant documents ( $\beta$ ).

# Rocchio Algorithm in Practice

- In practice, relevance feedback is very helpful to increase recall in situations where that is important.
  - This is also the most likely situation where a user is happy to provide feedback.
  - Positive feedback (i.e. identifying relevant documents) is more useful than negative feedback (i.e. identifying non-relevant documents), so in most systems  $\gamma < \beta$ .
  - Reasonable values might be  $\alpha = 1$ ,  $\beta = 0.75$  and  $\gamma = 0.15$ .
    - Some systems only support positive feedback, which is effectively setting  $\gamma = 0$ .

# Relevance Feedback Assumptions

- The user must have enough knowledge to make an **initial query that is close** to the relevant documents.
  - Without this, they will not be shown any relevant documents that they can give feedback on.
- Some problems that relevance feedback cannot solve by itself:
  - **Misspellings**: if the user's query does not spell terms correctly.
  - **Cross-language IR**: where documents in different languages will not be close in the vector space (although this depends on the vector representation).
  - **Vocabulary Mismatch**: where the user uses a different word/phrase to describe something (e.g. "laptop" versus "notebook computer").

# Relevance Feedback Assumptions

- Relevant documents are assumed to be **close to each other** in a cluster.
- **But** sometimes relevant documents can be in several clusters:
  - Documents might use different vocabulary (e.g. astronaut vs. cosmonaut).
  - A query that combines two very different things (e.g. Football players who studied at UCD).
  - General concepts that combine many more specific subjects (e.g. "felines" includes both wild animals and domestic pets).

# Pseudo Relevance Feedback

- In practice, users generally do not wish to provide feedback (known as **explicit** feedback).
- Instead, **pseudo relevance feedback** (also called **blind relevance feedback**) can be used: assume the top  $k$  ranked documents are relevant and then apply a relevance feedback algorithm.
- Quite successful in practice, although it can have problems in some situations.

# Indirect Relevance Feedback

- If users do not provide explicit feedback by actively marking documents as relevant or nonrelevant, other sources of evidence can be gathered to get **implicit** feedback. For example:
  - Which documents does the user choose to click on?
  - How long did the user spend viewing a document?
  - How long did the user spend browsing the results?



# Summary

- **Relevance feedback** allows the user to participate in retrieval by informing the system of results that are relevant and/or ones that are non-relevant.
  - **Example:** Rocchio algorithm based on vector representation.
- However, this is **not popular in practice**, so other ways are used to simulate feedback or get some feedback from other sources:
  - **Pseudo Relevance Feedback** (or “blind relevance feedback”: assume top  $k$  documents are relevant).
  - **Implicit Relevance Feedback** (evidence from other sources: clicks and user behaviour).